Elijah Campbell

campbell_salvus@gmail.com
(905) 244-6116
Ottawa, Canada
https://www.linkedin.com/in/elijah-campbell-16b513123/
https://github.com/DevDataMLOps

PROFESSIONAL SUMMARY

Data and Cloud Engineer focused on designing and operating scalable, cloud-native data platforms for analytics and SaaS. Proven track record building ETL and streaming pipelines, API-backed data services, and production-grade orchestration using Python, Apache Spark, Airflow, Kafka and AWS. Delivered platform improvements that cut end-to-end processing latency by 30% and reduced infrastructure spend by £10,000 per quarter. Skilled in performance tuning, data quality, and handover documentation for cross-functional teams.

SKILLS

Technical Skills

- Python
- SQL
- Apache Spark
- Apache Airflow
- Apache Kafka
- Apache Hadoop
- FastAPI
- AWS
- Dagster
- Databricks
- NumPy
- Google BigQuery
- PostgreSQL
- Docker
- Flask
- Terraform
- Pandas
- MongoDB
- Amazon Redshift
- Scala
- dbt (Data Build Tool)
- Kubernetes
- Data Validation
- Data Quality
- ETL Pipeline Design
- Stream Processing
- Cloud-Oriented Architecture
- CI/CD
- Prometheus
- Grafana
- Data Vault
-PySpark
- Snowflake

Core Skills

- System architecture and scalability
- ETL design and orchestration
- Performance optimization and tuning
- Stream processing and event-driven ingestion
- Data modeling and schema design
- Real-time analytics strategy
- Cost optimization and resource management
- CI/CD and reproducible deployments
- Technical documentation
- Cross-functional collaboration
- Agile delivery
- Data quality and validation

WORK EXPERIENCE

Data Engineer Internship - *Amdari* | May 2026 – Present · London, United Kingdom

- Designed and implemented scalable cloud data pipelines that reduced processing latency by 30% across client environments, using Spark on AWS EMR and Airflow orchestration
- Architect storage and warehousing optimizations that cut infrastructure costs by £10,000 per quarter through partitioning, tiered storage and query optimization
- Optimized database schemas and tuned queries to lower production query latency by 40% via indexing, partition strategies and query refactoring
- Enabled real-time analytics by implementing event-driven ingestion with Kafka and Kinesis, improving time-to-insight for operational dashboards by 3x

Data/Cloud Engineer – *10Alytics* | May 2026 – Oct 2026 · Alberta, Canada

- Increased data processing efficiency by 30% by developing scalable ETL pipelines using Python, SQL, and Apache Airflow, reducing daily pipeline execution time from 10 hours to 7 hours.
- Reduced PostgreSQL query latency by 40% through indexing, partitioning, and schema optimization, cutting average query response time from 5s to 3s.

- Saved £10,000 per quarter in cloud storage costs by redesigning PostgreSQL tables and implementing AWS S3 archival strategies, improving data retention and accessibility.
- Accelerated big data transformations by 60% using Apache Spark (PySpark), processing datasets over 1TB in under 45 minutes, supporting high-volume analytics.
- Implemented automated data validation scripts in Python and SQL, improving data quality scores by 25% and reducing reporting errors.

Lead Data Pipeline Engineer - *Project RISING (10Alytics Global Hackathon 2026)* | Jul 2026 – Present · Remote

- Led design and delivery of ETL pipelines for climate and public health datasets that validated, cleaned and processed 1,690 records for downstream analytics using Python
- Developed FastAPI-based REST services to expose health intelligence endpoints, reducing data retrieval time for consumers by 60%
- Contributed scalable system architecture and documentation for cloud-oriented processing workflows, enabling reproducible deployments and stakeholder handover

EDUCATION

Master Data Engineer (MDe) in Data/Cloud Engineer - *10Alytics* | 2026 (in progress)

PROJECTS

Project RISING (10Alytics Global Hackathon 2026) - *Python, FastAPI, Apache Spark, AWS*

Designed resilient ETL pipelines and REST APIs to process climate and public health datasets; validated, cleaned and processed 1,690 records and produced API and architecture documentation for downstream analytics.