

Elijah Campbell

campbell_salvus@gmail.com

(905)244-6116

Ottawa, Canada

<https://www.linkedin.com/in/elijah-campbell-16b513123/>

<https://github.com/DevDataMLOps>

PROFESSIONAL SUMMARY

Experienced Data and Cloud Engineer specializing in building scalable, cloud-native data platforms for SaaS and analytics products. Delivered pipeline and warehouse optimizations that reduced end-to-end processing latency by 30% and lowered infrastructure spend by £10,000 per quarter, applying Python, Apache Spark, Airflow and AWS at scale. Skilled in designing ETL/streaming architectures, API-backed data services with FastAPI, and implementing data quality and validation for production analytics.

SKILLS

Technical Skills

- | | | |
|------------------|-------------------|-------------------------------|
| - Python | - Google BigQuery | - Data Validation |
| - SQL | - PostgreSQL | - Data Quality |
| - Apache Spark | - Docker | - ETL Pipeline Design |
| - Apache Airflow | - Flask | - Stream Processing |
| - Apache Kafka | - Terraform | - Cloud-Oriented Architecture |
| - Apache Hadoop | - Pandas | - CI/CD |
| - FastAPI | - MongoDB | - Prometheus |
| - AWS | - Amazon Redshift | - Grafana |
| - Dagster | - Scala | - Data Vault |

Core Skills

- | | | |
|---------------------------------------|---|----------------------------------|
| - System architecture and scalability | - Data modeling and schema design | - Technical documentation |
| - ETL design and orchestration | - Real-time analytics strategy | - Cross-functional collaboration |
| - Performance optimization and tuning | - Cost optimization and resource management | - Agile delivery |

WORK EXPERIENCE

Data Engineer Internship - *Amdari* | May 2026 – Present · Ottawa, Canada

- Designed and implemented scalable cloud data pipelines that reduced processing latency by 30% across client environments, using Spark on AWS EMR and Airflow orchestration
- Architect storage and warehousing optimizations that cut infrastructure costs by £10,000 per quarter through partitioning, tiered storage and query optimization
- Optimized database schemas and tuned queries to lower production query latency by 40% via indexing, partition strategies and query refactoring
- Enabled real-time analytics by implementing event-driven ingestion with Kafka and Kinesis, improving time-to-insight for operational dashboards by 3x

Lead Data Pipeline Engineer - *Project RISING (10Alytics Global Hackathon 2026)* | Jul 2026 – Present · Remote

- Led design and delivery of ETL pipelines for climate and public health datasets that validated, cleaned and processed 1,690 records for downstream analytics using Python
- Developed FastAPI-based REST services to expose health intelligence endpoints, reducing data retrieval time for consumers by 60%

- Contributed scalable system architecture and documentation for cloud-oriented processing workflows, enabling reproducible deployments and stakeholder handover

EDUCATION

Master of Science (MSc) in Data/Cloud Engineer - *10Alytics* | 2026 (in progress)

PROJECTS

Project RISING (10Alytics Global Hackathon 2026) - *Python, FastAPI, Apache Spark, AWS*

Designed resilient ETL pipelines and REST APIs to process climate and public health datasets; validated, cleaned and processed 1,690 records and produced API and architecture documentation for downstream analytics.